



**19 MAT204 MATHEMATICS FOR
INTELLIGENT SYSTEMS - III
ADMM AND ITS APPLICATIONS**

GROUP 3

TEAM MEMBERS

Dinesh Kumar M R	- CB.EN.U4AIE20011
Rishi Karthigayan S	- CB.EN.U4AIE20057
N T Shrish Surya	- CB.EN.U4AIE20067
Tsaliki Satya Ganesh Kumar	- CB.EN.U4AIE20073

TABLE OF CONTENTS

Contents

Declaration.....	4
Acknowledgment.....	5
Abstract.....	6
Introduction:	7
Augmented Lagrangian and the Method of Multipliers:	8
ADMM and its various forms:	9
ADMM Form-1:	9
ADMM Form-2:	10
Applications of ADMM:	11
Linear Programming:.....	11
MATLAB Implementation :	12
Example for solving a LP:	13
Quadratic Programming:.....	14
MATLAB Implementation :	15
Example for solving a QP:.....	17
Basis Pursuit:.....	17
ADMM formulation:.....	18
MATLAB Implementation :	18
Example problem solved using ADMM:.....	20
L1 – Regularized Logistic Regression:	21
Matlab Implementation:.....	21
Example Problem using ADMM:	24
Support Vector Machine(SVM):	31
Matlab Implementation:.....	31
Example Problem using ADMM:	33
Intersection of Polyhedra:	37
Matlab Implementation:.....	38
Example Problem using ADMM:	39
Sparse Inverse Covariance Selection:	43
Matlab Implementation:.....	44
Example Problem using ADMM:	45
Inference:	49
Conclusion:	50

Declaration

We hereby declare that the project titled “ADMM and its applications” submitted to the Center for Excellence in Computational Engineering and Networking is a record of the original work done under the guidance of Dr. Neethu Mohan, Assistant Professor in Computational Engineering and Networking, Amrita University.

Signature of the faculty

Acknowledgment

The completion of this project has been made possible by the efforts of many. We would like to thank Dr. Sasangan Ramanathan, Dean - Engineering, and our Mathematics for Intelligent Systems-III faculty, Dr. Neethu Mohan, for their unstinting support. We would also like to extend our gratitude to our friends and family who have aided us throughout the project.

Abstract

Many problems of recent interest in statistics and machine learning can be posed in the framework of convex optimization. Due to the explosion in size and complexity of modern datasets, it is increasingly important to be able to solve problems with a very large number of features or training examples. As a result, both the decentralized collection or storage of these datasets as well as accompanying distributed solution methods are either necessary or at least highly desirable. In this review, we argue that the alternating direction method of multipliers is well suited to distributed convex optimization, and in particular to large-scale problems arising in statistics, machine learning, and related areas. The method was developed in the 1970s, with roots in the 1950s, and is equivalent or closely related to many other algorithms, such as dual decomposition, the method of multipliers, Douglas–Rachford splitting, Spingarn’s method of partial inverses, Dykstra’s alternating projections, Bregman iterative algorithms for ℓ_1 problems, proximal methods, and others. After briefly surveying the theory and history of the algorithm, we discuss applications to a wide variety of statistical and machine learning problems of recent interest, including the basis pursuit, regularised logistic regression, sparse inverse covariance selection, support vector machines, and many others. We also discuss linear programming, quadratic programming, and the intersection of polyhedra.

Introduction:

The alternating direction method of multipliers (ADMM) is an algorithm that solves convex optimization problems by breaking them into smaller pieces, each of which is then easier to handle. It has recently found wide application in a number of areas. ADMM is a simple and powerful iterative algorithm for convex optimization problems. It is much faster than conventional methods. It has emerged as a powerful technique for large-scale structured optimization.

ADMM extends the method of multipliers in such a way that we get back some of the decomposability (i.e. ability to parallelize) of standard dual ascent algorithms. It also gives us a flexible framework for incorporating many types of convex constraints. ADMM can be viewed as an attempt to blend the benefits of dual decomposition and augmented Lagrangian methods for constrained optimization.

Augmented Lagrangian and the Method of Multipliers:

Augmented lagrangian method was developed to bring robustness to the dual ascent method. The augmented lagrangian is denoted as $L_\rho(x, y)$. The augmented lagrangian for minimizing the function $f(x)$ subjected to $Ax = b$ is

$$L_\rho(x, y) = f(x) + \lambda(Ax - b) + (\rho/2) \|Ax - b\|_2^2$$

where $\rho > 0$, it is called penalty parameter and $(\rho/2) \|Ax - b\|_2^2$ is called a penalty term after modifying this problem is clearly equivalent to the original problem since for any feasible x the term added to the objective is zero.

Now we have to minimize the x value . for that we can apply dual ascent to our modified problem.

$$\begin{aligned} x^{k+1} &= L_\rho(x, z^k, \lambda^k) \\ z^{k+1} &= L_\rho(x^{k+1}, z, \lambda^k) \\ \lambda^{k+1} &= \lambda^k + \rho(Ax^{k+1} + Bz^{k+1} - c) \end{aligned}$$

At first we have to perform x -minimization and then z -minimization step. Thus they update alternatively

ADMM and its various forms:

ADMM Form-1:

Consider a need to minimize $f(x)+g(x)$,

ADMM works on the principle of increasing the number of variables and solving each of them via multiple iteration. Hence the problem can be re-formulated to,

$$\min f(x)+g(z)$$

$$\text{Subject to } x-z=0$$

The augmented Lagrangian can be written as :

$$L(x,y,z) = f(x)+g(z) + \langle y, x-z \rangle + (\rho/2) \cdot \|x - z\|_2^2$$

$(\rho/2) \cdot \|x - z\|_2^2$ is the augmented Lagrangian which is used to unconstrained the optimization and decrease the iterations to obtain the solution. While obtaining for nth iteration for a variable , the rest are treated as constants.

$$x^{k+1} = \arg \min_x \left(f(x) + \underbrace{\langle y^k, x \rangle}_{y^T x} + \frac{\rho}{2} \|x - z^k\|_2^2 \right)$$

In this iteration for x , the constant and other variables are neglected. Hence,

$$\begin{aligned}
x^{k+1} &= \arg \min_x \left(f(x) + \frac{\rho}{2} \left\| x - z^k + \frac{1}{\rho} y^k \right\|_2^2 \right) \\
z^{k+1} &= \arg \min_z \left(g(z) + \frac{\rho}{2} \left\| x^{k+1} - z + \frac{1}{\rho} y^k \right\|_2^2 \right) \\
y^{k+1} &= y^k + \rho (x^{k+1} - z^{k+1})
\end{aligned}$$

Substituting,

$$u^k = \frac{1}{\rho} y^k, \quad \lambda = \frac{1}{\rho}$$

$$\begin{aligned}
x^{k+1} &= \arg \min_x \left(f(x) + \frac{1}{2\lambda} \left\| x - (z^k - u^k) \right\|_2^2 \right) \\
z^{k+1} &= \arg \min_z \left(g(z) + \frac{1}{2\lambda} \left\| z - (x^{k+1} + u^k) \right\|_2^2 \right) \\
y^{k+1} &= y^k + \rho (x^{k+1} - z^{k+1})
\end{aligned}$$

ADMM Form-2:

Consider, $\min f(x)$

Subject to $x \in C$

This above problem is converted to $\min f(x) + g(z)$, Subject to $x-z=0$
Where $g(x)$ is an indicator function which takes the values 0 if $x \in C$ and infinity otherwise.

The augmented Lagrangian is written as ,

$$L(x, z, u) = f(x) + g(z) + \left(\frac{1}{2\lambda} \right) \|x - z + u\|_2^2$$

$$\begin{aligned}
x^{k+1} &= \arg \min_x \left(f(x) + \frac{1}{2\lambda} \left\| x - z^k + u^k \right\|_2^2 \right) \\
z^{k+1} &= \Pi_C (x^{k+1} + u^k) \\
u^{k+1} &= u^k + x^{k+1} - z^{k+1}
\end{aligned}$$

Where π_C represents Projection onto C.

Shrinkage function in a Nutshell :

When a threshold is fixed, any input above the threshold will be unaffected and if the input is less than the threshold it will be equated to 0.

Applications of ADMM:

Linear Programming:

In linear programming, the objective is to minimize a closed convex function

$$f(x) = \frac{1}{2} x^T P x + q^T x + r \quad ; \text{ Subject to } Ax=B$$

The problem will be converted to the form 2 of ADMM :

$$\min f(x) + g(z)$$

$$\text{Subject to } x - z = 0$$

Where $g(z)$ will be the indicator function.

$$L(x, z, u) = f(x) + g(z) + \left(\frac{1}{2}\lambda\right) \cdot \|x - z + u\|_2^2 \quad \text{[Lagrangian]}$$

Partial differentiation w.r.t x, z, u and substituting to 0

$$\frac{\partial L}{\partial x} = 0 \Rightarrow Px^{k+1} + q + \frac{1}{\lambda}(x^{k+1} - z^k + u^k) + A^T v = 0$$

$$\left(P + \frac{1}{\lambda}I\right)x^{k+1} + q + A^T v - \frac{1}{\lambda}(z^k - u^k) = 0 \quad \text{--(1)}$$

$$\frac{\partial L}{\partial v} = 0 \Rightarrow Ax^{k+1} - b = 0 \quad \text{--(2)}$$

$$\begin{bmatrix} P + \rho \cdot I & A^T \\ A & 0 \end{bmatrix} \cdot \begin{bmatrix} x^{k+1} \\ u \end{bmatrix} = \begin{bmatrix} \rho(z^k - u^k) - q \\ b \end{bmatrix}$$

$$\begin{bmatrix} P + \rho \cdot I & A^T \\ A & 0 \end{bmatrix} \cdot \begin{bmatrix} x^{k+1} \\ v \end{bmatrix} = \begin{bmatrix} P + \rho \cdot I & A^T \\ A & 0 \end{bmatrix}^{-1} \cdot \begin{bmatrix} \rho(z^k - u^k) - q \\ b \end{bmatrix} \quad | \text{EQ-1}$$

-- (From 1 and 2)

where $\rho = 1/\lambda$

The very same methodology can be implemented for solving Quadratic equations on x when subject to $Ax=b$. In case of a subject of inequality, the methodology will differ. Usually solving for a linear program, P will be a 0 matrix which reduces $f(x)$ to $c'(x)+r$;

MATLAB Implementation :

```
function [z, history] = linprog(c, A, b, rho, alpha)
t_start = tic;

%Global Constants and Defaults

QUIET      = 0;
MAX_ITER   = 1000;
ABSTOL     = 1e-4;
RELTOL     = 1e-2;

%Data preprocessing
[m n] = size(A);

%ADMM Solver
x = zeros(n,1);
z = zeros(n,1);
u = zeros(n,1);
```

```

if ~QUIET
    fprintf('%3s\t%10s\t%10s\t%10s\t%10s\t%10s\n', 'iter', ...
        'r norm', 'eps pri', 's norm', 'eps dual', 'objective');
end

for k = 1:MAX_ITER

    % x-update
    tmp = [ rho*eye(n), A'; A, zeros(m) ] \ [ rho*(z - u) - c; b ];
    x = tmp(1:n);

    % z-update with relaxation
    zold = z;
    x_hat = alpha*x + (1 - alpha)*zold;
    z = pos(x_hat + u);

    u = u + (x_hat - z);

    % diagnostics, reporting, termination checks

    history.objval(k) = objective(c, x);

    history.r_norm(k) = norm(x - z);
    history.s_norm(k) = norm(-rho*(z - zold));

    history.eps_pri(k) = sqrt(n)*ABSTOL + RELTOL*max(norm(x), norm(-z));
    history.eps_dual(k) = sqrt(n)*ABSTOL + RELTOL*norm(rho*u);

    if ~QUIET
        fprintf('%3d\t%10.4f\t%10.4f\t%10.4f\t%10.4f\t%10.2f\n', k, ...
            history.r_norm(k), history.eps_pri(k), ...
            history.s_norm(k), history.eps_dual(k), history.objval(k));
    end

    if (history.r_norm(k) < history.eps_pri(k) && ...
        history.s_norm(k) < history.eps_dual(k))
        break;
    end
end

if ~QUIET
    toc(t_start);
end

end

function obj = objective(c, x)
    obj = c'*x;
end

```

Example for solving a LP:

```

randn('state', 0);
rand('state', 0);

n = 5; % dimension of x

```

```

m = 4; % number of equality constraints

c = rand(n,1) + 0.5; % create nonnegative price vector with mean 1
x0 = abs(randn(n,1)); % create random solution vector

A = abs(randn(m,n)); % create random, nonnegative matrix A
b = A*x0;

[x history] = linprog(c, A, b, 1.0, 1.0);

```

Output:

iter	r norm	eps pri	s norm	eps dual	objective
1	0.0000	0.0209	2.0673	0.0002	3.91
2	0.0824	0.0227	0.4471	0.0010	3.64
3	0.4063	0.0247	0.1864	0.0051	3.45
4	0.1817	0.0233	0.1293	0.0069	3.58
5	0.0107	0.0223	0.1107	0.0068	3.70
6	0.0505	0.0221	0.0229	0.0063	3.72
7	0.0225	0.0222	0.0161	0.0061	3.70
8	0.0014	0.0223	0.0137	0.0061	3.69
9	0.0063	0.0224	0.0028	0.0062	3.69

Elapsed time is 0.014902 seconds.

Quadratic Programming:

For solving Quadratic equations Subject to lower and upper bound.

$$f(x) = \frac{1}{2}x^T \cdot P \cdot x + q^T x + r \quad ; \text{ Subject to : } lb < x < ub$$

The problem will be converted to the form 2 of ADMM :

$$\min f(x) + g(z)$$

$$\text{Subject to } x - z = 0$$

where $g(z)$ will be the indicator function.

$$L(x, z, u) = f(x) + g(z) + (\rho/2) \cdot ||x - z + u||_2^2 \quad | \text{Lagrangian}$$

From EQ-1 : When A, b and v are null –

$$x = [P + \rho.I]^{-1} * [\rho.(z - u) - q]$$

| I is an Identity Matrix

For increasing efficiency and speed of computation , Consider R to be an upper triangular matrix obtained from Cholesky Factorisation of $[P+p.I]$.

Cholesky Factorisation :

When A is a Hermitian and Positive Deficient Matrix on factorisation we get R, Where $A = R * R^T$

Hence to increase efficiency ,

$$x^{k+1} = R * R^{-1} * [\rho.(z^k - u^k) - q]$$

| Where $[P + \rho.I]^{-1} = R * R^T$

$$z = \min(ub, \max(lb, x^{k+1} + u^k))$$

$$u^{k+1} = u^k + (x^{k+1} - z)$$

MATLAB Implementation :

```
function [z, history] = quadprog(P, q, r, lb, ub, rho, alpha)

t_start = tic;

%Global constants and defaults
QUIET      = 0;
MAX_ITER   = 1000;
ABSTOL     = 1e-4;
RELTOL     = 1e-2;

%Data preprocessing
n = size(P,1);

%ADMM solver
x = zeros(n,1);
z = zeros(n,1);
u = zeros(n,1);

if ~QUIET
    fprintf('%3s\t%10s\t%10s\t%10s\t%10s\t%10s\n', 'iter', ...
        'r norm', 'eps pri', 's norm', 'eps dual', 'objective');
end

for k = 1:MAX_ITER
```

```

if k > 1
    x = R \ (R' \ (rho*(z - u) - q));
else
    R = chol(P + rho*eye(n));
    x = R \ (R' \ (rho*(z - u) - q));
end

% z-update with relaxation
zold = z;
x_hat = alpha*x + (1-alpha)*zold;
z = min(ub, max(lb, x_hat + u));

% u-update
u = u + (x_hat - z);

% diagnostics, reporting, termination checks
history.objval(k) = objective(P, q, r, x);

history.r_norm(k) = norm(x - z);
history.s_norm(k) = norm(-rho*(z - zold));

history.eps_pri(k) = sqrt(n)*ABSTOL + RELTOL*max(norm(x), norm(-z));
history.eps_dual(k) = sqrt(n)*ABSTOL + RELTOL*norm(rho*u);

if ~QUIET
    fprintf('%3d\t%10.4f\t%10.4f\t%10.4f\t%10.4f\t%10.2f\n', k, ...
            history.r_norm(k), history.eps_pri(k), ...
            history.s_norm(k), history.eps_dual(k), history.objval(k));
end

if (history.r_norm(k) < history.eps_pri(k) && ...
    history.s_norm(k) < history.eps_dual(k))
    break;
end
end

if ~QUIET
    toc(t_start);
end

end

function obj = objective(P, q, r, x)
    obj = 0.5*x'*P*x + q'*x + r;
end

```


Example for solving a QP:

```
randn('state', 0);
rand('state', 0);

n = 100;

% generate a well-conditioned positive definite matrix
% (for faster convergence)
P = rand(n);
P = P + P';
[V D] = eig(P);
P = V*diag(1+rand(n,1))*V';

q = randn(n,1);
r = randn(1);

l = randn(n,1);
u = randn(n,1);
lb = min(l,u);
ub = max(l,u);

%Solve problem
[x history] = quadprog(P, q, r, lb, ub, 1.0, 1.0);
```

Output:

iter	r norm	eps pri	s norm	eps dual	objective
1	5.4585	0.0508	4.9772	0.0556	-23.62
2	2.0050	0.0546	0.8151	0.0734	-10.06
3	1.2096	0.0558	0.2823	0.0845	-4.36
4	0.7197	0.0562	0.1188	0.0912	-0.39
5	0.4287	0.0564	0.0578	0.0952	2.14
6	0.2564	0.0564	0.0331	0.0976	3.70
7	0.1541	0.0564	0.0199	0.0990	4.64
8	0.0931	0.0564	0.0120	0.0999	5.21
9	0.0566	0.0564	0.0073	0.1004	5.55
10	0.0345	0.0564	0.0044	0.1007	5.75

Elapsed time is 0.067481 seconds.

Basis Pursuit:

The objective is very simple to minimise the values satisfying $Ax = b$

minimize $\|x\|_1$ subject to $Ax = b$

ADMM formulation:

minimize $f(x)g(x)$ subject to x - z vector

where,

$$f(x) = \|x\|_1$$

$$g(z) = \begin{cases} 0, & Az = b \\ \infty & \text{otherwise} \end{cases}$$

$g(z)$ is like a penalty function when it satisfies the condition $Az=b$, it will be 0 otherwise it will shoot up to a large value.

The ADMM Algorithm will be :

$$x^{k+1} := \Pi(z^k - u^k)$$

$$z^{k+1} := S_{1/\rho}(x^{k+1} + u^k)$$

$$u^{k+1} := u^k + x^{k+1} - z^{k+1},$$

where Π is projection onto $\{x \in \mathbb{R}^n \mid Ax = b\}$.

The x -update, which involves solving a linearly-constrained minimum Euclidean norm problem, can be written explicitly as

$$x^{k+1} := (I - A^T(AA^T)^{-1}A)(z^k - u^k) + A^T(AA^T)^{-1}b.$$

Now, let $P = (I - A^T(AA^T)^{-1}A)$ and $q = A^T(AA^T)^{-1}b$.

$$\text{Hence, } x^{k+1} = P*(z^k - u^k) + q$$

MATLAB Implementation :

```
function [z, history] = basis_pursuit(A, b, rho, alpha)
```

```
t_start = tic;
```

```
QUIET      = 0;  
MAX_ITER   = 1000;  
ABSTOL     = 1e-4;  
RELTOL     = 1e-2;
```

```

[m n] = size(A);

x = zeros(n,1);
z = zeros(n,1);
u = zeros(n,1);

if ~QUIET
    fprintf('%3s\t%10s\t%10s\t%10s\t%10s\n', 'iter', ...
        'r norm', 'eps pri', 's norm', 'eps dual', 'objective');
end

% precompute static variables for x-update (projection on to Ax=b)
AA_t = A*A';
P = eye(n) - A' * (AA_t \ A);
q = A' * (AA_t \ b);

for k = 1:MAX_ITER
    % x-update
    x = P*(z - u) + q;

    % z-update with relaxation
    zold = z;
    x_hat = alpha*x + (1 - alpha)*zold;
    z = shrinkage(x_hat + u, 1/rho);

    u = u + (x_hat - z);

    % diagnostics, reporting, termination checks
    history.objval(k) = objective(A, b, x);

    history.r_norm(k) = norm(x - z);
    history.s_norm(k) = norm(-rho*(z - zold));

    history.eps_pri(k) = sqrt(n)*ABSTOL + RELTOL*max(norm(x), norm(-z));
    history.eps_dual(k) = sqrt(n)*ABSTOL + RELTOL*norm(rho*u);

    if ~QUIET
        fprintf('%3d\t%10.4f\t%10.4f\t%10.4f\t%10.4f\t%10.2f\n', k, ...
            history.r_norm(k), history.eps_pri(k), ...
            history.s_norm(k), history.eps_dual(k), history.objval(k));
    end

    if (history.r_norm(k) < history.eps_pri(k) && ...
        history.s_norm(k) < history.eps_dual(k))
        break;
    end
end
end

```

```

if ~QUIET
    toc(t_start);
end

end

function obj = objective(A, b, x)
    obj = norm(x,1);
end

function y = shrinkage(a, kappa)
    y = max(0, a-kappa) - max(0, -a-kappa);
end

```

Example problem solved using ADMM:

```

rand('seed', 0);
randn('seed', 0);

n = 3;
m = 1;
A = randn(m,n);

x = sprandn(n, 1, 0.1*n);
b = A*x;
xtrue = x;
[x history] = basis_pursuit(A, b, 1.0, 1.0);

```

Output:

iter	r norm	eps pri	s norm	eps dual	objective
1	0.3091	0.0033	0.0000	0.0033	0.44
2	0.3091	0.0033	0.0000	0.0064	0.44
3	0.3091	0.0033	0.0000	0.0094	0.44
4	0.2362	0.0033	0.0871	0.0118	0.44
5	0.0743	0.0034	0.2242	0.0122	0.39
6	0.0771	0.0042	0.0820	0.0118	0.48
7	0.0518	0.0040	0.0137	0.0115	0.44
8	0.0060	0.0037	0.0248	0.0115	0.36
9	0.0090	0.0036	0.0082	0.0115	0.36
10	0.0055	0.0037	0.0019	0.0116	0.35
11	0.0004	0.0037	0.0027	0.0116	0.35

Elapsed time is 0.035173 seconds.

L1 – Regularized Logistic Regression:

Logistic regression is widely used in machine learning for classification problems. It is well-known that regularization is required to avoid over-fitting, especially when there is only a small number of training examples, or when there are a large number of parameters to be learned. In particular, L1 regularized logistic regression is often used for feature selection and has been shown to have good generalization performance in the presence of many irrelevant features.

L1 regularized logistic regression is a workhorse of machine learning: it is widely used for many classification problems, particularly ones with many features. L1 regularized logistic regression requires solving a convex optimization problem. However, a drawback is that standard algorithms for solving convex optimization problems do not scale well enough to handle the large datasets encountered in many practical settings.

Matlab Implementation:

```
function [z, history] = logreg(A, b, mu, rho, alpha)
t_start = tic;
```

Global constants and defaults

```
QUIET = 0;
MAX_ITER = 1000;
ABSTOL = 1e-4;
RELTOL = 1e-2;
```

Data preprocessing

```
[m, n] = size(A);
```

ADMM solver

```
x = zeros(n+1,1);
```

```

z = zeros(n+1,1);

u = zeros(n+1,1);

if ~QUIET

    fprintf('%3s\t%10s\t%10s\t%10s\t%10s\t%10s\n', 'iter', ...
        'r norm', 'eps pri', 's norm', 'eps dual', 'objective');

end

for k = 1:MAX_ITER

    % x-update

    x = update_x(A, b, u, z, rho);

    % z-update with relaxation

    zold = z;

    x_hat = alpha*x + (1-alpha)*zold;

    z = x_hat + u;

    z(2:end) = shrinkage(z(2:end), (m*mu)/rho);

    u = u + (x_hat - z);

    % diagnostics, reporting, termination checks

    history.objval(k) = objective(A, b, mu, x, z);

    history.r_norm(k) = norm(x - z);

    history.s_norm(k) = norm(rho*(z - zold));

    history.eps_pri(k) = sqrt(n)*ABSTOL + RELTOL*max(norm(x), norm(z));

    history.eps_dual(k) = sqrt(n)*ABSTOL + RELTOL*norm(rho*u);

    if ~QUIET

        fprintf('%3d\t%10.4f\t%10.4f\t%10.4f\t%10.4f\t%10.2f\n', k, ...
            history.r_norm(k), history.eps_pri(k), ...
            history.s_norm(k), history.eps_dual(k), history.objval(k));

    end

    if history.r_norm(k) < history.eps_pri(k) && ...

```

```

        history.s_norm(k) < history.eps_dual(k)

        break;

    end

end

if ~QUIET

    toc(t_start);

end

end

function obj = objective(A, b, mu, x, z)

    m = size(A,1);

    obj = sum(log(1 + exp(-A*x(2:end) - b*x(1)))) + m*mu*norm(z,1);

end

function x = update_x(A, b, u, z, rho, x0)

    % solve the x update

    % minimize [ -logistic(x_i) + (rho/2)||x_i - z^k + u^k||^2 ]

    % via Newton's method; for a single subsystem only.

    alpha = 0.1;

    BETA = 0.5;

    TOLERANCE = 1e-5;

    MAX_ITER = 50;

    [m n] = size(A);

    I = eye(n+1);

    if exist('x0', 'var')

        x = x0;

    else

        x = zeros(n+1,1);

    end

```

```

C = [-b -A];

f = @(w) (sum(log(1 + exp(C*w))) + (rho/2)*norm(w - z + u).^2);

for iter = 1:MAX_ITER

    fx = f(x);

    g = C'*(exp(C*x)./(1 + exp(C*x))) + rho*(x - z + u);

    H = C' * diag(exp(C*x)./(1 + exp(C*x)).^2) * C + rho*I;

    dx = -H\g;    % Newton step

    dfx = g'*dx; % Newton decrement

    if abs(dfx) < TOLERANCE

        break;

    end

    % backtracking

    t = 1;

    while f(x + t*dx) > fx + alpha*t*dfx

        t = BETA*t;

    end

    x = x + t*dx;

end

end

function z = shrinkage(a, kappa)

    z = max(0, a-kappa) - max(0, -a-kappa);

end

```

Example Problem using ADMM:

Generate problem data

```

rand('seed', 0);

randn('seed', 0);

n = 5;

```



```

m = 20;

w = sprandn(n, 1, 0.1); % N(0,1), 10% sparse
v = randn(1);          % random intercept
X = sprandn(m, n, 10/n);
btrue = sign(X*w + v);

% noise is function of problem size use 0.1 for large problem
b = sign(X*w + v + sqrt(0.1)*randn(m,1)); % labels with noise
A = spdiags(b, 0, m, m) * X;
ratio = sum(b == 1)/(m);

mu = 0.1 * 1/m * norm((1-ratio)*sum(A(b==1,:),1) + ratio*sum(A(b==-1,:),1),
'inf');

x_true = [v; w];

```

Solve problem

```
[x history] = logreg(A, b, mu, 1.0, 1.0);
```

iter	r norm	eps pri	s norm	eps dual	objective
1	0.3774	0.0199	1.8188	0.0040	5.56
2	0.0416	0.0276	0.9767	0.0042	4.78
3	0.0202	0.0340	0.6855	0.0041	4.32
4	0.0472	0.0390	0.5355	0.0038	4.06
5	0.0515	0.0431	0.4411	0.0036	3.89
6	0.0486	0.0466	0.3754	0.0034	3.78
7	0.0435	0.0497	0.3265	0.0033	3.70
8	0.0380	0.0524	0.2886	0.0033	3.64
9	0.0328	0.0549	0.2581	0.0033	3.60
10	0.0283	0.0571	0.2331	0.0033	3.57
11	0.0565	0.0591	0.2047	0.0030	3.55
12	0.0262	0.0609	0.1815	0.0029	3.54

13	0.0142	0.0625	0.1636	0.0029	3.53
14	0.0106	0.0640	0.1487	0.0029	3.53
15	0.0089	0.0653	0.1360	0.0028	3.52
16	0.0076	0.0666	0.1250	0.0028	3.52
17	0.0066	0.0677	0.1152	0.0028	3.52
18	0.0058	0.0688	0.1067	0.0028	3.52
19	0.0051	0.0698	0.0990	0.0028	3.53
20	0.0046	0.0707	0.0921	0.0028	3.53
21	0.0041	0.0715	0.0859	0.0028	3.53
22	0.0037	0.0723	0.0803	0.0028	3.53
23	0.0033	0.0731	0.0753	0.0028	3.54
24	0.0030	0.0738	0.0706	0.0028	3.54
25	0.0027	0.0744	0.0664	0.0028	3.55
26	0.0025	0.0751	0.0625	0.0028	3.55
27	0.0023	0.0756	0.0589	0.0028	3.55
28	0.0021	0.0762	0.0556	0.0028	3.56
29	0.0019	0.0767	0.0525	0.0028	3.56
30	0.0018	0.0772	0.0496	0.0028	3.56
31	0.0016	0.0777	0.0470	0.0028	3.57
32	0.0015	0.0781	0.0445	0.0028	3.57
33	0.0014	0.0785	0.0422	0.0028	3.58
34	0.0013	0.0789	0.0401	0.0028	3.58
35	0.0012	0.0793	0.0381	0.0028	3.58
36	0.0011	0.0797	0.0362	0.0028	3.59
37	0.0011	0.0800	0.0344	0.0028	3.59
38	0.0010	0.0804	0.0327	0.0028	3.59
39	0.0009	0.0807	0.0312	0.0028	3.59
40	0.0009	0.0810	0.0297	0.0028	3.60

41	0.0008	0.0812	0.0283	0.0028	3.60
42	0.0008	0.0815	0.0270	0.0028	3.60
43	0.0007	0.0818	0.0257	0.0028	3.61
44	0.0007	0.0820	0.0246	0.0028	3.61
45	0.0006	0.0823	0.0234	0.0028	3.61
46	0.0006	0.0825	0.0224	0.0028	3.61
47	0.0006	0.0827	0.0214	0.0028	3.61
48	0.0005	0.0829	0.0204	0.0028	3.62
49	0.0005	0.0831	0.0196	0.0028	3.62
50	0.0005	0.0833	0.0187	0.0028	3.62
51	0.0005	0.0835	0.0179	0.0028	3.62
52	0.0004	0.0836	0.0171	0.0028	3.62
53	0.0004	0.0838	0.0164	0.0028	3.63
54	0.0004	0.0839	0.0157	0.0028	3.63
55	0.0004	0.0841	0.0150	0.0028	3.63
56	0.0004	0.0842	0.0144	0.0028	3.63
57	0.0003	0.0844	0.0138	0.0028	3.63
58	0.0003	0.0845	0.0132	0.0028	3.63
59	0.0003	0.0846	0.0126	0.0028	3.64
60	0.0003	0.0848	0.0121	0.0028	3.64
61	0.0003	0.0849	0.0116	0.0028	3.64
62	0.0003	0.0850	0.0111	0.0028	3.64
63	0.0003	0.0851	0.0107	0.0028	3.64
64	0.0002	0.0852	0.0102	0.0028	3.64
65	0.0002	0.0853	0.0098	0.0028	3.64
66	0.0002	0.0854	0.0094	0.0028	3.64
67	0.0002	0.0855	0.0090	0.0028	3.65
68	0.0002	0.0856	0.0087	0.0028	3.65

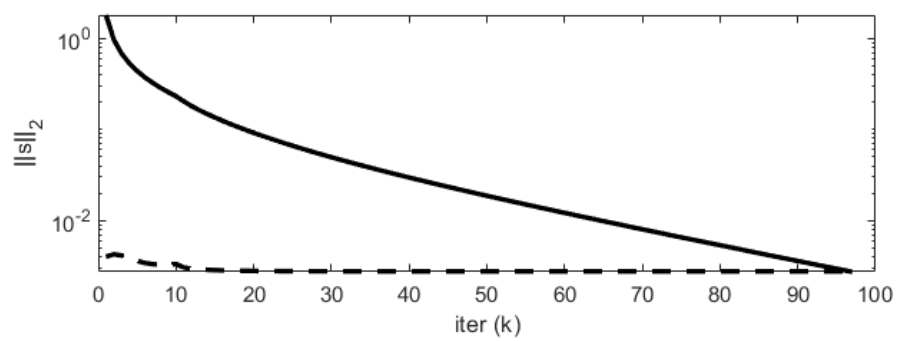
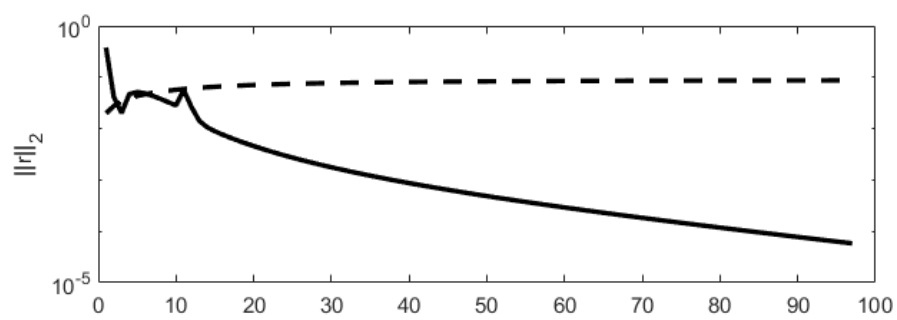
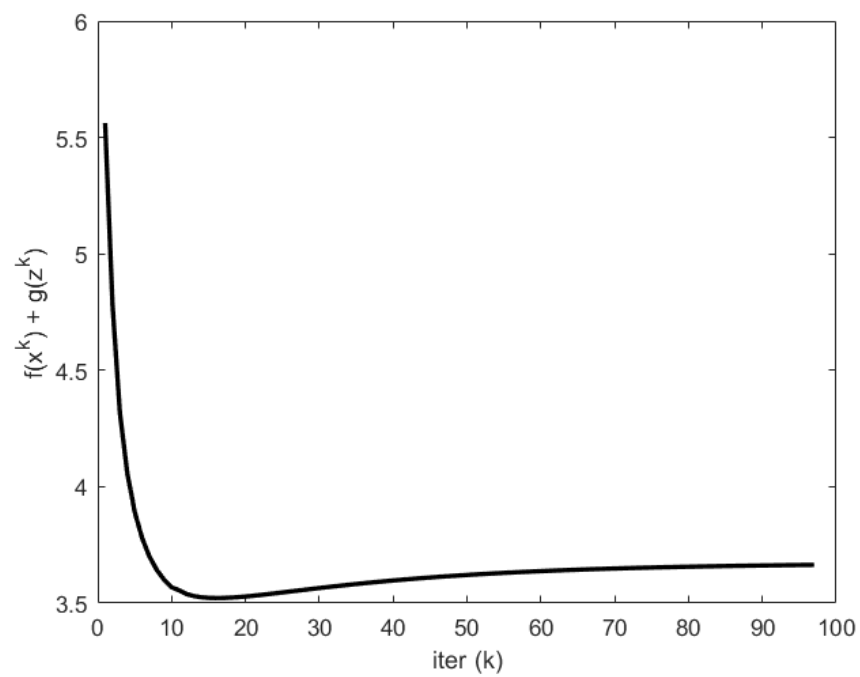
69	0.0002	0.0856	0.0083	0.0028	3.65
70	0.0002	0.0857	0.0080	0.0028	3.65
71	0.0002	0.0858	0.0077	0.0028	3.65
72	0.0002	0.0859	0.0074	0.0028	3.65
73	0.0002	0.0859	0.0071	0.0028	3.65
74	0.0002	0.0860	0.0068	0.0028	3.65
75	0.0001	0.0861	0.0065	0.0028	3.65
76	0.0001	0.0861	0.0063	0.0028	3.65
77	0.0001	0.0862	0.0060	0.0028	3.65
78	0.0001	0.0863	0.0058	0.0028	3.66
79	0.0001	0.0863	0.0056	0.0028	3.66
80	0.0001	0.0864	0.0053	0.0028	3.66
81	0.0001	0.0864	0.0051	0.0028	3.66
82	0.0001	0.0865	0.0049	0.0028	3.66
83	0.0001	0.0865	0.0047	0.0028	3.66
84	0.0001	0.0866	0.0046	0.0028	3.66
85	0.0001	0.0866	0.0044	0.0028	3.66
86	0.0001	0.0866	0.0042	0.0028	3.66
87	0.0001	0.0867	0.0040	0.0028	3.66
88	0.0001	0.0867	0.0039	0.0028	3.66
89	0.0001	0.0868	0.0037	0.0028	3.66
90	0.0001	0.0868	0.0036	0.0028	3.66
91	0.0001	0.0868	0.0034	0.0028	3.66
92	0.0001	0.0869	0.0033	0.0028	3.66
93	0.0001	0.0869	0.0032	0.0028	3.66
94	0.0001	0.0869	0.0031	0.0028	3.66
95	0.0001	0.0870	0.0029	0.0028	3.66
96	0.0001	0.0870	0.0028	0.0028	3.66

97	0.0001	0.0870	0.0027	0.0028	3.66
----	--------	--------	--------	--------	------

Elapsed time is 0.258555 seconds.

Reporting

```
K = length(history.objval);
h = figure;
plot(1:K, history.objval, 'k', 'MarkerSize', 10, 'LineWidth', 2);
ylabel('f(x^k) + g(z^k)'); xlabel('iter (k)');
g = figure;
subplot(2,1,1);
semilogy(1:K, max(1e-8, history.r_norm), 'k', ...
    1:K, history.eps_pri, 'k--', 'LineWidth', 2);
ylabel('||r||_2');
subplot(2,1,2);
semilogy(1:K, max(1e-8, history.s_norm), 'k', ...
    1:K, history.eps_dual, 'k--', 'LineWidth', 2);
ylabel('||s||_2'); xlabel('iter (k)');
```



Support Vector Machine(SVM):

Support Vector Machine(SVM) is a supervised machine learning algorithm used for both classification and regression. Though we say regression problems it is best suited for classification also. The objective of the SVM algorithm is to find a hyperplane in an N-dimensional space that distinctly classifies the data points. The dimension of the hyperplane depends upon the number of features. If the number of input features is two, then the hyperplane is just a line. If the number of input features is three, then the hyperplane becomes a 2-D plane. It becomes difficult when the number of features exceeds three.

Matlab Implementation:

```
function [xave, history] = linear_svm(A, lambda, p, rho, alpha)

t_start = tic;
```

Global constants and defaults

```
QUIET      = 0;
MAX_ITER   = 1000;
ABSTOL     = 1e-4;
RELTOL     = 1e-2;
```

Data preprocessing

```
[m, n] = size(A);
N = max(p);
% group samples together
for i = 1:N,
    tmp{i} = A(p==i,:);
end
A = tmp;
```

ADMM solver

```
x = zeros(n,N);
z = zeros(n,N);
u = zeros(n,N);

if ~QUIET
    fprintf('%3s\t%10s\t%10s\t%10s\t%10s\t%10s\n', 'iter', ...
        'r norm', 'eps pri', 's norm', 'eps dual', 'objective');
```

```

end

for k = 1:MAX_ITER

    % x-update
    for i = 1:N,
        cvx_begin quiet
            variable x_var(n)
            minimize ( sum(pos(A{i}*x_var + 1)) + rho/2*sum_square(x_var -
z(:,i) + u(:,i)) )
            cvx_end
            x(:,i) = x_var;
        end
        xave = mean(x,2);

    % z-update with relaxation
    zold = z;
    x_hat = alpha*x + (1-alpha)*zold;
    z = N*rho/(1/lambda + N*rho)*mean( x_hat + u, 2 );
    z = z*ones(1,N);

    % u-update
    u = u + (x_hat - z);

    % diagnostics, reporting, termination checks
    history.objval(k) = objective(A, lambda, p, x, z);

    history.r_norm(k) = norm(x - z);
    history.s_norm(k) = norm(-rho*(z - zold));

    history.eps_pri(k) = sqrt(n)*ABSTOL + RELTOL*max(norm(x), norm(-z));
    history.eps_dual(k) = sqrt(n)*ABSTOL + RELTOL*norm(rho*u);

    if ~QUIET
        fprintf('%3d\t%10.4f\t%10.4f\t%10.4f\t%10.4f\t%10.2f\n', k, ...
            history.r_norm(k), history.eps_pri(k), ...
            history.s_norm(k), history.eps_dual(k), history.objval(k));
    end

    if (history.r_norm(k) < history.eps_pri(k) && ...
        history.s_norm(k) < history.eps_dual(k))
        break;
    end
end

if ~QUIET
    toc(t_start);
end
end

function obj = objective(A, lambda, p, x, z)
    obj = hinge_loss(A,x) + 1/(2*lambda)*sum_square(z(:,1));
end

function val = hinge_loss(A,x)
    val = 0;
    for i = 1:length(A)
        val = val + sum(pos(A{i}*x(:,i) + 1));
    end
end
end

```


Example Problem using ADMM:

```
rand('seed', 0);
randn('seed', 0);

n = 2;
m = 200;
N = m/2;
M = m/2;

% positive examples
Y = [1.5+0.9*randn(1,0.6*N), 1.5+0.7*randn(1,0.4*N);
     2*(randn(1,0.6*N)+1), 2*(randn(1,0.4*N)-1)];

% negative examples
X = [-1.5+0.9*randn(1,0.6*M), -1.5+0.7*randn(1,0.4*M);
     2*(randn(1,0.6*M)-1), 2*(randn(1,0.4*M)+1)];

x = [X Y];
y = [ones(1,N) -ones(1,M)];
A = [ -(ones(n,1)*y).*x' -y'];
xdat = x';
lambda = 1.0;

% partition the examples up in the worst possible way
% (subsystems only have positive or negative examples)
p = zeros(1,m);
p(y == 1) = sort(randi([1 10], sum(y==1),1));
p(y == -1) = sort(randi([11 20], sum(y==-1),1));
```

Solve problem

```
[x history] = linear_svm(A, lambda, p, 1.0, 1.0);
```

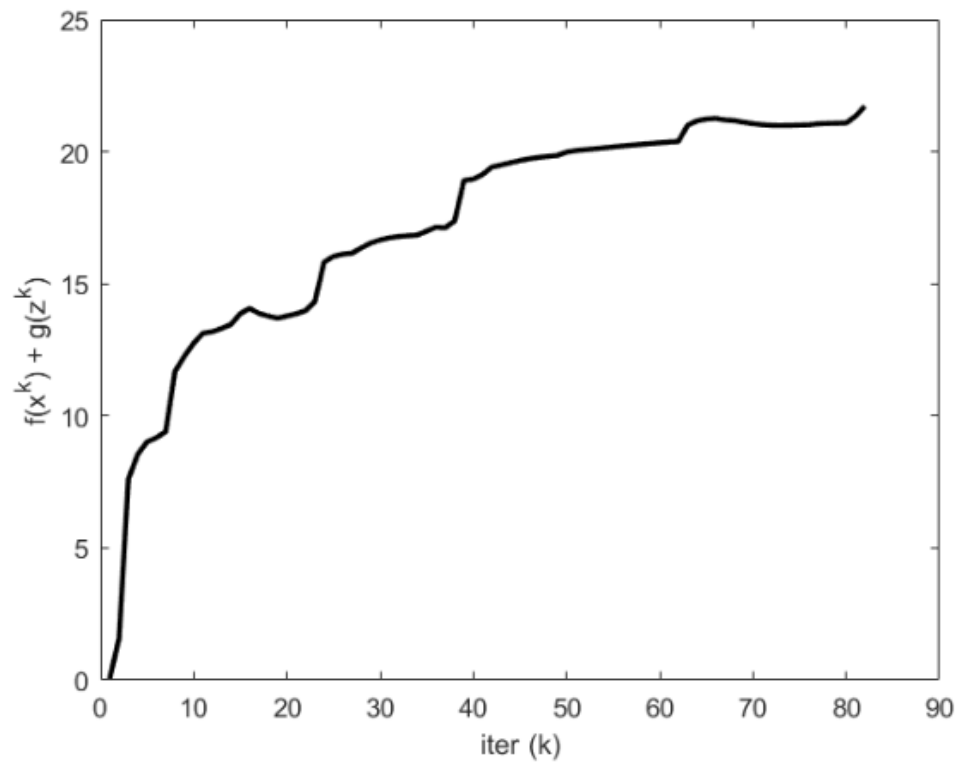
iter	r norm	eps pri	s norm	eps dual	objective
1	2.9530	0.0297	1.4676	0.0297	0.05
2	1.9801	0.0331	1.6816	0.0488	1.58
3	1.0197	0.0458	1.3087	0.0532	7.60
4	0.9436	0.0561	1.0746	0.0554	8.54
5	0.9444	0.0649	0.9008	0.0562	9.01
6	0.9496	0.0719	0.7339	0.0560	9.17
7	0.9143	0.0770	0.5603	0.0605	9.40
8	0.7525	0.0804	0.3914	0.0668	11.68
9	0.7373	0.0821	0.2032	0.0728	12.25
10	0.6569	0.0823	0.0943	0.0782	12.75
11	0.6171	0.0816	0.1587	0.0832	13.13
12	0.5971	0.0803	0.1992	0.0881	13.18
13	0.5683	0.0783	0.2423	0.0929	13.31
14	0.5399	0.0759	0.2734	0.0975	13.46
15	0.4846	0.0735	0.2743	0.1016	13.87
16	0.4656	0.0712	0.2565	0.1054	14.07
17	0.4655	0.0691	0.2169	0.1092	13.88
18	0.4664	0.0675	0.1637	0.1131	13.78
19	0.4666	0.0663	0.1247	0.1171	13.69
20	0.4604	0.0655	0.1159	0.1210	13.78
21	0.4584	0.0650	0.1244	0.1251	13.86
22	0.4558	0.0646	0.1200	0.1291	13.98
23	0.4295	0.0644	0.0946	0.1330	14.32
24	0.3584	0.0644	0.0765	0.1358	15.82
25	0.3566	0.0646	0.0512	0.1385	16.04
26	0.3570	0.0649	0.0467	0.1411	16.12
27	0.3523	0.0653	0.0613	0.1438	16.16
28	0.3316	0.0658	0.0627	0.1464	16.36
29	0.3251	0.0663	0.0694	0.1490	16.54
30	0.3242	0.0669	0.0708	0.1514	16.66
31	0.3236	0.0675	0.0686	0.1539	16.74
32	0.3229	0.0682	0.0649	0.1564	16.79
33	0.3222	0.0687	0.0596	0.1588	16.82
34	0.3217	0.0692	0.0522	0.1614	16.84
35	0.3124	0.0697	0.0457	0.1638	16.99
36	0.3055	0.0701	0.0513	0.1662	17.14
37	0.3044	0.0706	0.0537	0.1686	17.12
38	0.2834	0.0711	0.0566	0.1708	17.37
39	0.2130	0.0716	0.0769	0.1723	18.91
40	0.2058	0.0722	0.0844	0.1737	18.96
41	0.1941	0.0728	0.0810	0.1750	19.13
42	0.1833	0.0735	0.0798	0.1762	19.42
43	0.1779	0.0742	0.0842	0.1773	19.50
44	0.1749	0.0749	0.0842	0.1785	19.58
45	0.1741	0.0756	0.0790	0.1796	19.66
46	0.1735	0.0763	0.0716	0.1807	19.73
47	0.1726	0.0769	0.0641	0.1817	19.78
48	0.1714	0.0774	0.0566	0.1828	19.82
49	0.1706	0.0779	0.0483	0.1839	19.85
50	0.1602	0.0783	0.0459	0.1849	19.98
51	0.1571	0.0787	0.0453	0.1859	20.04
52	0.1563	0.0791	0.0442	0.1869	20.07
53	0.1560	0.0795	0.0433	0.1878	20.11
54	0.1560	0.0798	0.0426	0.1887	20.14
55	0.1562	0.0801	0.0416	0.1897	20.18
56	0.1567	0.0804	0.0398	0.1906	20.21
57	0.1572	0.0807	0.0369	0.1915	20.25

58	0.1578	0.0810	0.0334	0.1924	20.28
59	0.1583	0.0812	0.0295	0.1933	20.31
60	0.1586	0.0814	0.0254	0.1942	20.34
61	0.1587	0.0816	0.0215	0.1950	20.36
62	0.1587	0.0818	0.0181	0.1959	20.39
63	0.1318	0.0819	0.0216	0.1965	21.00
64	0.1285	0.0820	0.0270	0.1971	21.17
65	0.1276	0.0821	0.0242	0.1976	21.23
66	0.1263	0.0821	0.0146	0.1981	21.26
67	0.1235	0.0821	0.0057	0.1986	21.20
68	0.1226	0.0820	0.0124	0.1991	21.18
69	0.1219	0.0820	0.0168	0.1997	21.11
70	0.1217	0.0819	0.0166	0.2003	21.06
71	0.1219	0.0819	0.0130	0.2010	21.02
72	0.1221	0.0818	0.0085	0.2016	21.00
73	0.1221	0.0817	0.0067	0.2023	20.99
74	0.1219	0.0817	0.0079	0.2029	21.00
75	0.1218	0.0816	0.0092	0.2036	21.01
76	0.1217	0.0815	0.0095	0.2043	21.01
77	0.1139	0.0815	0.0136	0.2049	21.05
78	0.1126	0.0815	0.0183	0.2055	21.07
79	0.1110	0.0815	0.0219	0.2061	21.08
80	0.1092	0.0816	0.0261	0.2068	21.09
81	0.0872	0.0817	0.0325	0.2073	21.34
82	0.0669	0.0817	0.0405	0.2076	21.72

Elapsed time is 779.912628 seconds.

Reporting

```
K = length(history.objval);
h = figure;
plot(1:K, history.objval, 'k', 'MarkerSize', 10, 'LineWidth', 2);
ylabel('f(x^k) + g(z^k)'); xlabel('iter (k)');
```

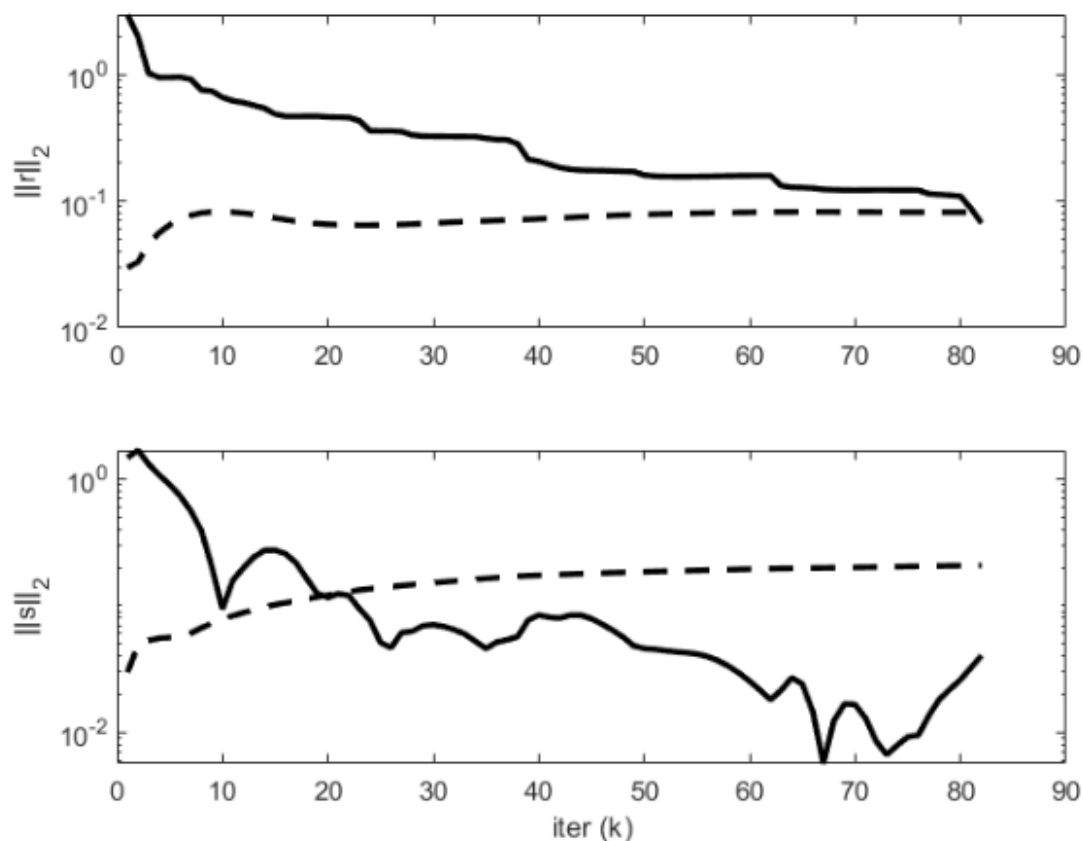


```

g = figure;
subplot(2,1,1);
semilogy(1:K, max(1e-8, history.r_norm), 'k', ...
    1:K, history.eps_pri, 'k--', 'LineWidth', 2);
ylabel('||r||_2');

subplot(2,1,2);
semilogy(1:K, max(1e-8, history.s_norm), 'k', ...
    1:K, history.eps_dual, 'k--', 'LineWidth', 2);
ylabel('||s||_2'); xlabel('iter (k)');

```



Intersection of Polyhedra:

The intersection detection problem is a subset of collision detection for more complex forms, which plays a major role in robotics, computer animation, and mechanical simulation.

Intersection detection is also used in computer graphics as a component of accelerated data structures such as bounding volume hierarchies and Kd-trees. In the latter scenario, an item of interest (such as a ray, a view frustum, or another hierarchy) is compared to the geometric forms that bind each node in the hierarchy to see if they overlap. Simple forms (boxes, spheres) for which quick intersection detection algorithms exist are used as bounding shapes.

Matlab Implementation:

```
function [z, history] = polyhedra_intersection(A1, b1, A2, b2, rho, alpha)

t_start = tic;
```

Global constants and defaults

```
QUIET      = 0;
MAX_ITER   = 1000;
ABSTOL     = 1e-4;
RELTOL     = 1e-2;

n = size(A1,2);
```

ADMM solver

```
x = zeros(n,1);
z = zeros(n,1);
u = zeros(n,1);

if ~QUIET
    fprintf('%3s\t%10s\t%10s\t%10s\t%10s\t%10s\n', 'iter', ...
        'r norm', 'eps pri', 's norm', 'eps dual', 'objective');
end

for k = 1:MAX_ITER

    % x-update
    % use cvx to find point in first polyhedra
    cvx_begin quiet
        variable x(n)
        minimize (sum_square(x - (z - u)))
        subject to
            A1*x <= b1
    cvx_end

    % z-update with relaxation
    zold = z;
    x_hat = alpha*x + (1 - alpha)*zold;
    % use cvx to find point in second polyhedra
    cvx_begin quiet
        variable z(n)
        minimize (sum_square(x_hat - (z - u)))
        subject to
            A2*z <= b2
    cvx_end

    u = u + (x_hat - z);

    % diagnostics, reporting, termination checks

    history.objval(k) = 0;
    history.r_norm(k) = norm(x - z);
    history.s_norm(k) = norm(-rho*(z - zold));

    history.eps_pri(k) = sqrt(n)*ABSTOL + RELTOL*max(norm(x), norm(-z));
    history.eps_dual(k) = sqrt(n)*ABSTOL + RELTOL*norm(rho*u);
```

```

    if ~QUIET
        fprintf('%3d\t%10.4f\t%10.4f\t%10.4f\t%10.4f\t%10.2f\n', k, ...
            history.r_norm(k), history.eps_pri(k), ...
            history.s_norm(k), history.eps_dual(k), history.objval(k));
    end

    if (history.r_norm(k) < history.eps_pri(k) && ...
        history.s_norm(k) < history.eps_dual(k))
        break;
    end
end

if ~QUIET
    toc(t_start);
end
end

```

Example Problem using ADMM:

```

randn('state', 0);
rand('state', 0);

n = 5;          % dimension of variable
m1 = 16;        % number of faces for polyhedra 1
m2 = 18;        % number of faces for polyhedra 2

c1 = 10*randn(n,1);      % center of polyhedra 1
c2 = -10*randn(n,1);     % center of polyhedra 2

% consider the following picture:
%
%      a1
% c -----> x
%
% from the center "c", we travel along vector "a1" (not necessarily a unit
% vector) until we reach x. at "x", a1'x = b. a point y is to the left of x
% if a1'y <= b.
%

% pick m1 random directions with different magnitudes
A1 = diag(1 + rand(m1,1))*randn(m1,n);
% the value of b is found by traveling from the center along the normal
% vectors in A1 and taking its inner product with A1.
b1 = diag(A1*(c1*ones(1,m1) + A1')));

% pick m2 random directions with different magnitudes
A2 = diag(1 + rand(m2,1))*randn(m2,n);
% the value of b is found by traveling from the center along the normal
% vectors in A1 and taking its inner product with A1.
b2 = diag(A2*(c2*ones(1,m2) + A2')));

% find the distance between the two polyhedra--make sure they overlap by
% checking if the distance is 0
cvx_begin quiet
    variables x(n) y(n)
    minimize sum_square(x - y)
    subject to

```

```

        A1*x <= b1
        A2*y <= b2
cvx_end

% if the distance is not 0, expand A1 and A2 by a little more than half the
% distance
if norm(x-y) > 1e-4,
    A1 = (1 + 0.5*norm(x-y))*A1;
    A2 = (1 + 0.5*norm(x-y))*A2;
    % recompute b's as appropriate
    b1 = diag(A1*(c1*ones(1,m1) + A1')) ;
    b2 = diag(A2*(c2*ones(1,m2) + A2')) ;
end

```

Solve problem

```
[x history] = polyhedra_intersection(A1, b1, A2, b2, 1.0, 1.0);
```

iter	r norm	eps pri	s norm	eps dual	objective
1	2.6152	0.1501	14.9907	0.0264	0.00
2	1.1149	0.1612	1.1554	0.0152	0.00
3	1.1944	0.1618	0.1560	0.0033	0.00
4	0.3059	0.1600	0.0000	0.0002	0.00
5	0.0000	0.1595	0.0000	0.0002	0.00

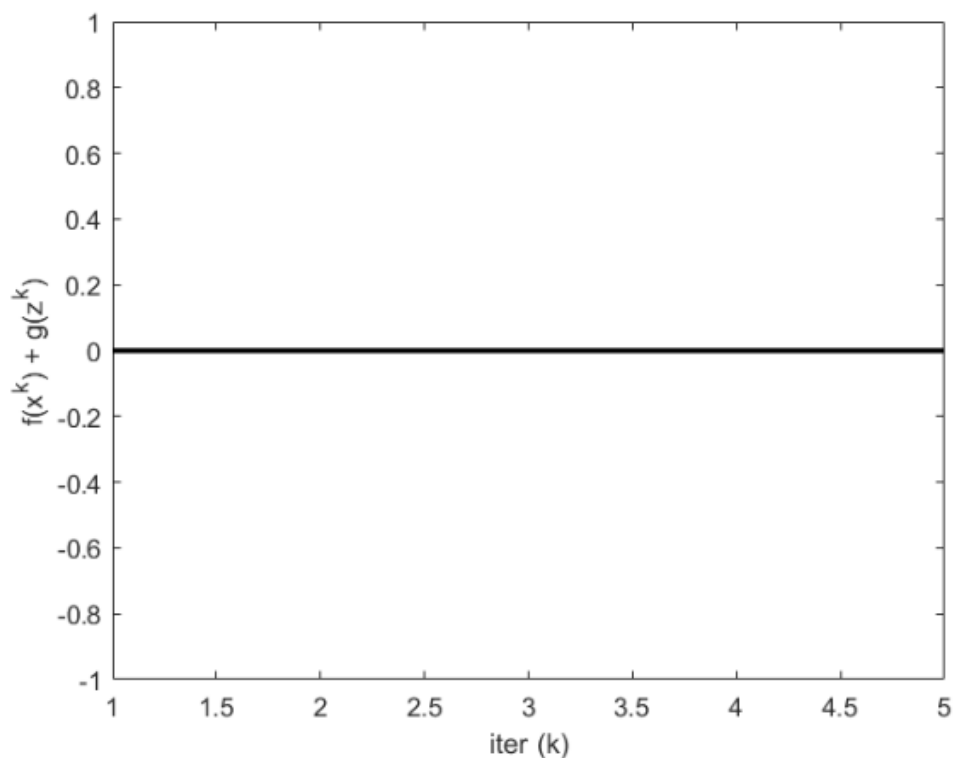
Elapsed time is 5.307592 seconds.

Reporting

```

K = length(history.objval);
h = figure;
plot(1:K, history.objval, 'k', 'MarkerSize', 10, 'LineWidth', 2);
ylabel('f(x^k) + g(z^k)'); xlabel('iter (k)');

```

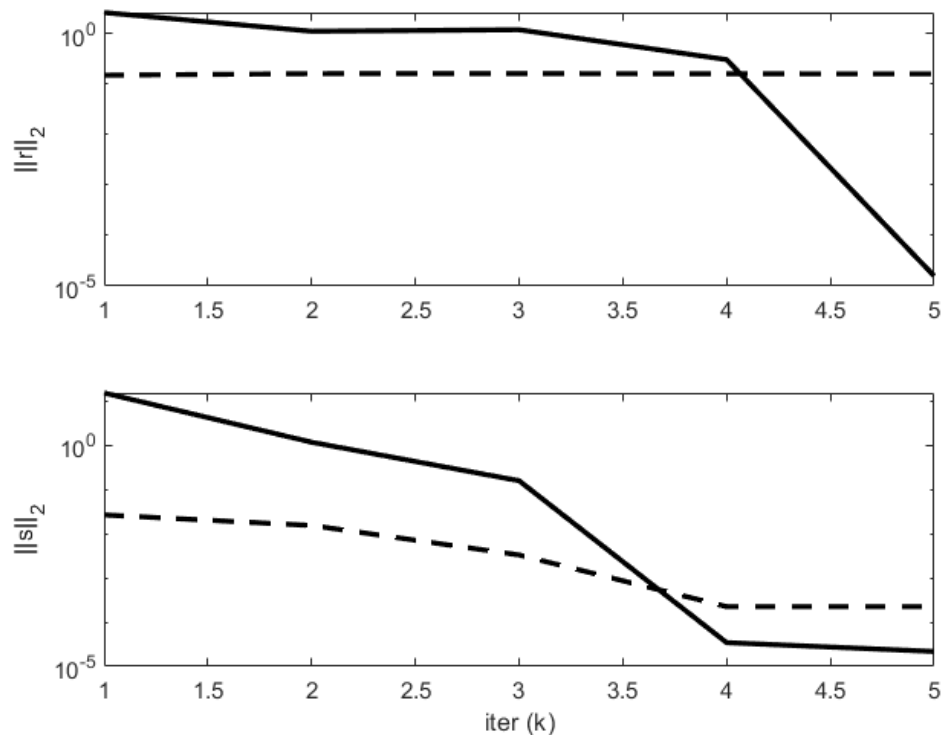



```

g = figure;
subplot(2,1,1);
semilogy(1:K, max(1e-8, history.r_norm), 'k', ...
    1:K, history.eps_pri, 'k--', 'LineWidth', 2);
ylabel('||r||_2');

subplot(2,1,2);
semilogy(1:K, max(1e-8, history.s_norm), 'k', ...
    1:K, history.eps_dual, 'k--', 'LineWidth', 2);
ylabel('||s||_2'); xlabel('iter (k)');

```



Compare to alternating projections

```

% Compare to alternating projections

MAX_ITER = 10;

x = zeros(n,1);
z = zeros(n,1);

for k = 1:MAX_ITER
    % x-update
    % use cvx to find point in first polyhedra
    cvx_begin quiet
        variable x(n)

```

```

    minimize (sum_square(x - z))

    subject to

        A1*x <= b1

cvx_end

% z-update with relaxation

zold = z;

% use cvx to find point in second polyhedra

cvx_begin quiet

    variable z(n)

    minimize (sum_square(x - z))

    subject to

        A2*z <= b2

cvx_end

history1.r_norm(k) = norm(x - z);

history1.s_norm(k) = norm((z - zold));

end

g = figure;

subplot(2,1,1);

semilogy(1:MAX_ITER, max(1e-8, history1.r_norm), 'k', 1:K, max(1e-8,
history.r_norm), 'r');

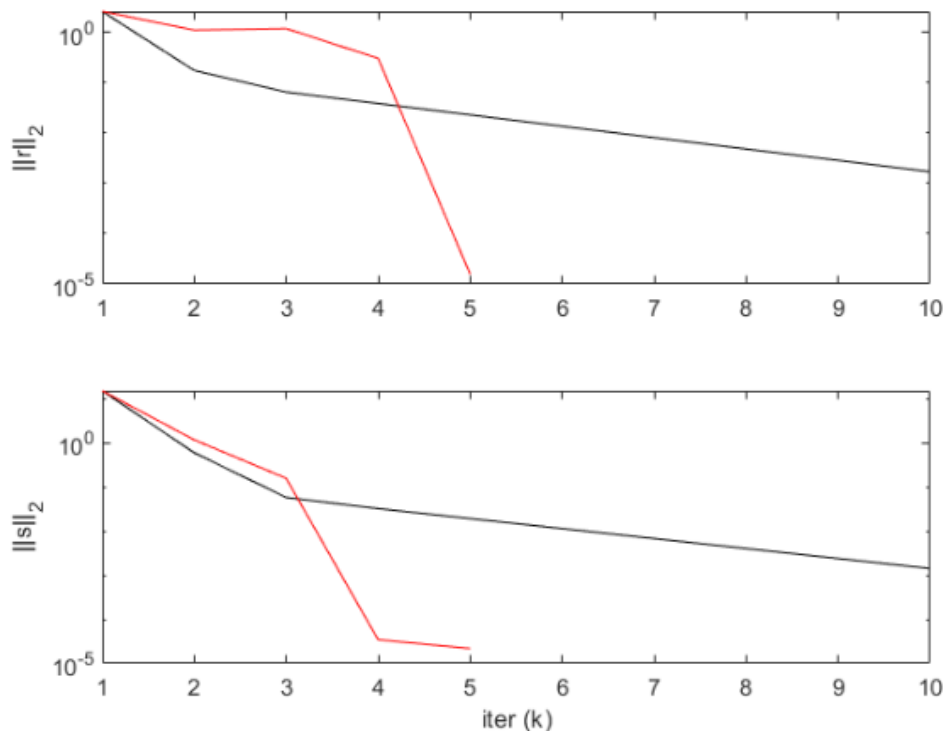
ylabel('||r||_2');

subplot(2,1,2);

semilogy(1:MAX_ITER, max(1e-8, history1.s_norm), 'k', 1:K, max(1e-8,
history.s_norm), 'r');

ylabel('||s||_2'); xlabel('iter (k)');

```



Sparse Inverse Covariance Selection:

It is a sequence of positive definite inverse covariance matrices that converge to a sparse matrix, while the sequence of Y^k 's is a sequence of sparse matrices that converges to a positive definite inverse covariance matrix. To estimate a probabilistic model (e.g. a Gaussian model), estimating the precision matrix, that is the inverse covariance matrix, is as important as estimating the covariance matrix. Indeed a Gaussian model is parametrized by the precision matrix. Given an empirical covariance matrix S ,

$$S = \frac{1}{N} \sum_{i=1}^N (x_i - \mu)(x_i - \mu)^T.$$

we find a sparse inverse covariance matrix P to represent the data.

$$\min_{P \succ 0} F(P) \stackrel{\text{def}}{=} L(P) + \lambda \|\text{vec}(P)\|_1, \quad L(P) = -\log \det(P) + \text{trace}(SP).$$

where L is the negative log-likelihood function and the l_1 term is a sparsity inducing regularizer.

Matlab Implementation:

```
function [Z, history] = covsel(D, lambda, rho, alpha)
```

```
t_start = tic;
```

Global constants and defaults

```
QUIET      = 0;  
MAX_ITER   = 1000;  
ABSTOL     = 1e-4;  
RELTOL     = 1e-2;
```

Data preprocessing

```
S = cov(D);  
n = size(S,1);
```

ADMM solver

```
X = zeros(n);  
Z = zeros(n);  
U = zeros(n);  
  
if ~QUIET  
    fprintf('%3s\t%10s\t%10s\t%10s\t%10s\t%10s\n', 'iter', ...  
           'r norm', 'eps pri', 's norm', 'eps dual', 'objective');  
end  
  
for k = 1:MAX_ITER  
  
    % x-update  
    [Q,L] = eig(rho*(Z - U) - S);  
    es = diag(L);  
    xi = (es + sqrt(es.^2 + 4*rho))./(2*rho);  
    X = Q*diag(xi)*Q';  
  
    % z-update with relaxation  
    Zold = Z;  
    X_hat = alpha*X + (1 - alpha)*Zold;  
    Z = shrinkage(X_hat + U, lambda/rho);  
  
    U = U + (X_hat - Z);  
  
    % diagnostics, reporting, termination checks  
  
    history.objval(k) = objective(S, X, Z, lambda);  
  
    history.r_norm(k) = norm(X - Z, 'fro');  
    history.s_norm(k) = norm(-rho*(Z - Zold), 'fro');  
  
    history.eps_pri(k) = sqrt(n*n)*ABSTOL + RELTOL*max(norm(X, 'fro'),  
norm(Z, 'fro'));  
    history.eps_dual(k) = sqrt(n*n)*ABSTOL + RELTOL*norm(rho*U, 'fro');  
  
    if ~QUIET  
        fprintf('%3d\t%10.4f\t%10.4f\t%10.4f\t%10.4f\t%10.2f\n', k, ...  
                history.r_norm(k), history.eps_pri(k), ...
```

```

        history.s_norm(k), history.eps_dual(k), history.objval(k));
    end

    if (history.r_norm(k) < history.eps_pri(k) && ...
        history.s_norm(k) < history.eps_dual(k))
        break;
    end
end

if ~QUIET
    toc(t_start);
end
end

function obj = objective(S, X, Z, lambda)
    obj = trace(S*X) - log(det(X)) + lambda*norm(Z(:), 1);
end

function y = shrinkage(a, kappa)
    y = max(0, a-kappa) - max(0, -a-kappa);
end

```

Example Problem using ADMM:

```

randn('seed', 0);
rand('seed', 0);

n = 100;    % number of features
N = 10*n;   % number of samples

% generate a sparse positive definite inverse covariance matrix
Sinv      = diag(abs(ones(n,1)));
idx       = randsample(n^2, 0.001*n^2);
Sinv(idx) = ones(numel(idx), 1);
Sinv = Sinv + Sinv';    % make symmetric
if min(eig(Sinv)) < 0    % make positive definite
    Sinv = Sinv + 1.1*abs(min(eig(Sinv)))*eye(n);
end
S = inv(Sinv);

% generate Gaussian samples
D = mvnrnd(zeros(1,n), S, N);

```

Solve problem

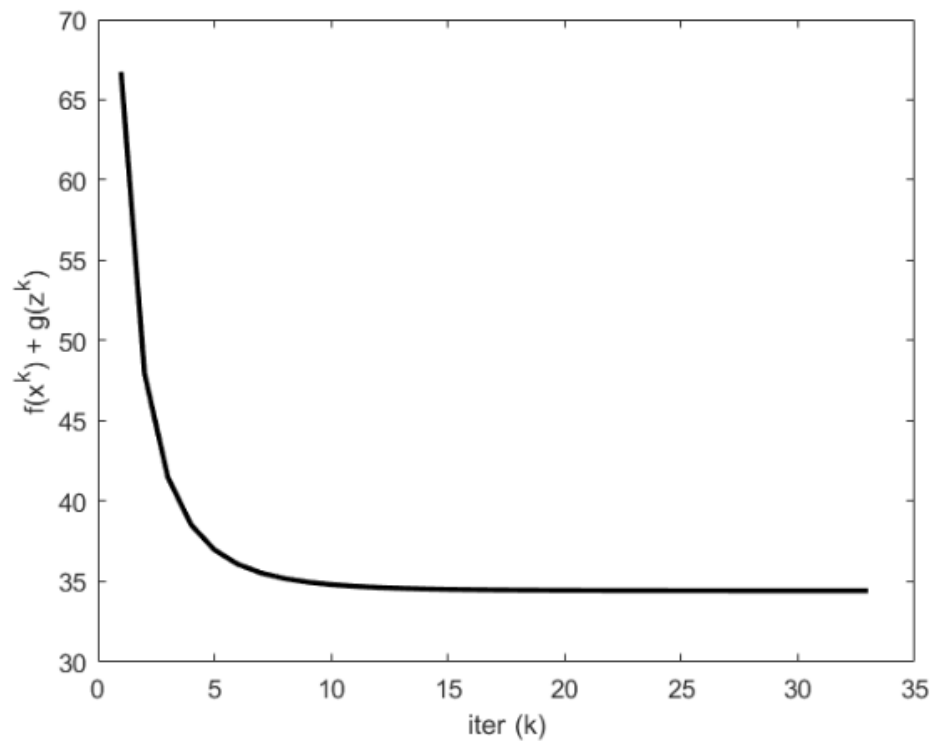
```
[X, history] = covsel(D, 0.01, 1, 1);
```

iter	r norm	eps pri	s norm	eps dual	objective
1	0.5604	0.0878	7.6521	0.0156	66.74
2	0.2677	0.1230	3.6651	0.0176	47.97
3	0.0890	0.1453	2.2727	0.0181	41.51
4	0.0341	0.1608	1.5731	0.0182	38.54
5	0.0223	0.1720	1.1563	0.0183	36.98
6	0.0193	0.1806	0.8831	0.0183	36.08
7	0.0168	0.1872	0.6927	0.0183	35.54
8	0.0147	0.1924	0.5544	0.0183	35.19
9	0.0126	0.1967	0.4507	0.0183	34.97
10	0.0110	0.2001	0.3710	0.0183	34.82
11	0.0092	0.2029	0.3087	0.0183	34.71
12	0.0079	0.2053	0.2592	0.0183	34.64
13	0.0074	0.2072	0.2192	0.0183	34.59
14	0.0060	0.2088	0.1867	0.0183	34.55
15	0.0051	0.2102	0.1600	0.0183	34.52
16	0.0044	0.2114	0.1378	0.0183	34.50
17	0.0036	0.2124	0.1192	0.0183	34.48
18	0.0031	0.2132	0.1036	0.0183	34.47
19	0.0028	0.2140	0.0904	0.0183	34.46
20	0.0025	0.2146	0.0792	0.0182	34.46
21	0.0023	0.2151	0.0696	0.0182	34.45
22	0.0019	0.2156	0.0613	0.0182	34.45
23	0.0016	0.2160	0.0542	0.0182	34.44
24	0.0014	0.2163	0.0480	0.0182	34.44
25	0.0013	0.2166	0.0426	0.0182	34.44
26	0.0011	0.2169	0.0379	0.0182	34.44
27	0.0012	0.2171	0.0337	0.0182	34.44
28	0.0040	0.2173	0.0299	0.0182	34.44
29	0.0021	0.2175	0.0265	0.0182	34.44
30	0.0013	0.2177	0.0236	0.0182	34.43
31	0.0009	0.2178	0.0211	0.0182	34.43
32	0.0008	0.2179	0.0188	0.0182	34.43
33	0.0007	0.2181	0.0169	0.0182	34.43

Elapsed time is 0.256118 seconds.

Reporting

```
K = length(history.objval);  
X_admm = X;  
  
h = figure;  
plot(1:K, history.objval, 'k', 'MarkerSize', 10, 'LineWidth', 2);  
ylabel('f(x^k) + g(z^k)'); xlabel('iter (k)');
```

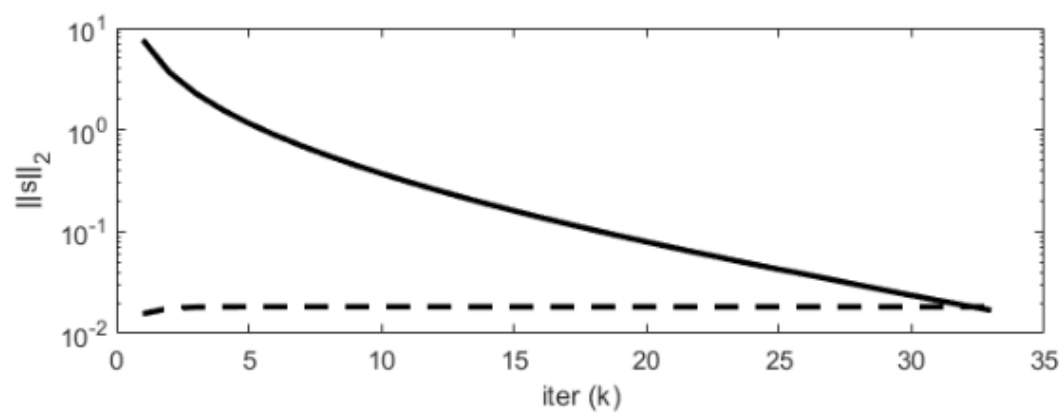
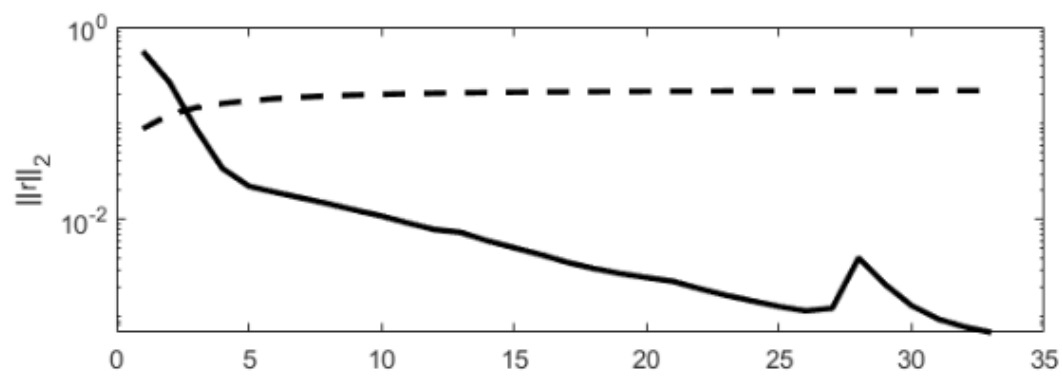


```

g = figure;
subplot(2,1,1);
semilogy(1:K, max(1e-8, history.r_norm), 'k', ...
    1:K, history.eps_pri, 'k--', 'LineWidth', 2);
ylabel('||r||_2');

subplot(2,1,2);
semilogy(1:K, max(1e-8, history.s_norm), 'k', ...
    1:K, history.eps_dual, 'k--', 'LineWidth', 2);
ylabel('||s||_2'); xlabel('iter (k)');

```



Inference:

- ADMM has the benefit of being extremely simple to implement, and it maps onto several standard distributed programming models reasonably well.
- ADMM was developed over a generation ago, with its roots stretching far in advance of the Internet, distributed and cloud computing systems, massive high-dimensional datasets, and the associated large scale applied statistical problems.
- Despite this, it appears to be well suited to the modern regime and has the important benefit of being quite general in its scope and applicability.
- ADMM builds on existing algorithms for single machines, and so can be viewed as a modular coordination algorithm that ‘incentivizes’ a set of simpler algorithms to collaborate to solve much larger global problems together than they could on their own.

Conclusion: